Microway, Inc.
12 Richards Road
Plymouth, MA 02360
Phone: 508.746.7341
Fax: 508.746.4678
http://www.microway.com

QUOTATION

Quote # MWYQ23588
Date: 4/9/2018
Microway Confidential

Quote To:

Pomona College
Asya Shklyar
333 N. College Way
Claremont, CA 91711

(909) 607-9853
asya_shklyar@hotmail.com

Ship To:

Pomona College
Asya Shklyar
333 N. College Way
Claremont, CA 91711

(909) 607-9853
asya_shklyar@hotmail.com

Sales Rep.

Samantha Wheeler
508-732-5526
swheeler@microway.com

GPU-Accelerated Workstation for Deep Learning

Qty	Description	Unit Price	Extended Price

Microway is a small business, woman owned and operated. We are building many clusters at any one time and have built thousands of custom clusters for universities, government research labs and agencies, and corporations. Microway has been in the scientific computing business since 1982.

MWYQ23350-01

1 **DGX Station Deep Learning Workstation with Tesla V100** $46,775 $46,775

NVIDIA DGX Station Deep Learning Workstation with "Volta" V100 32GB GPUs
Purpose-built by NVIDIA for Deep Learning
Arrives with all hardware and software fully-integrated

One 20-core Intel Xeon E5-2698v4 CPU
256GB DDR4 2133MHz System Memory
Four Tesla V100 32GB "Volta" GPU Accelerators with NVLink 2.0
(a total of 20,480 NVIDIA CUDA cores; 2,560 Tensor cores)
Four 1.92TB SSD (one for OS and three in RAID0 for high-speed cache)
Dual X540 10GbE Ethernet ports (10GBase-T RJ45 ports)
Ubuntu Desktop Linux Operating System

Software stack includes:
* DIGITS training system
* NVIDIA Deep Learning SDK with latest versions of CUDA & cuDNN
* NVIDIA Cluster Portal (cloud or onsite)
* Online application repository with the major deep learning frameworks
* NVDocker containerized app deployment
* Managed app container creation and deployment
* Multi-Node management with telemetry, monitoring and alerts

Quiet, liquid-cooled tower form factor (<35 dB for office use)
Power Consumption: 1500W at full load (for standard office electrical outlet)

https://www.nvidia.com/en-us/data-center/dgx-station

Support Services for NVIDIA DGX Station "Volta" Deep Learning Workstation with V100 32GB (1-Year Comm/Govt)
Includes one year of 24x7 support and hardware warranty with 1-day SLA (replacement parts shipped)
Includes DGX System Software, Driver and Firmware Updates

For complete details, please see:
http://www.nvidia.com/object/dgx1-support.html

MWYQ23494

Qty	Description	Unit Price	Extended Price
1	**Microway 1U Xeon + Tesla GPU Server with NVLink**	$39,752	$39,752

NumberSmasher Dual Intel Xeon 1U 4-GPU NVLink 2.0 Server with
2000W High-Efficiency "Titanium" Redundant 1+1 Power Supplies
(power supplies require 208V power for redundancy under full load)
Two Intel Xeon Scalable Family processors (Socket P - up to 165W)
Three Intel UPI interconnects between CPU sockets (up to 10.4 GT/s)
Twelve slots for up to 1.5TB ECC DDR4-2666 memory
Dual Integrated Intel X540 10 Gigabit Ethernet ports (RJ45, 10GBase-T)
Integrated Intel C621 chipset and ASpeed AST2500 Graphics Controller
IPMI 2.0 w/ Virtual Media, KVM and Dedicated LAN Support
Four SXM2 sockets for NVIDIA NVLink 2.0-enabled GPUs
Up to 300GB/s Bi-directional Bandwidth between GPUs (6 links per GPU)
Four PCI-Express 3.0 x16 slots:
* two full-height, half-length slots (one via PLX; one via CPU)
* two low-profile, half-length slots (one via PLX; one via CPU)
Integrated SATA3 6Gbps Controller
Two hot-swap 2.5" SAS/SATA drive bays
Two internal 2.5" SAS/SATA drive bays
Internal ports: two SATA SuperDOM, one M.2 slot (PCIe 3.0 x4 NVMe via PCH), one TPM header, one serial COM header
Rear ports: two 10G LAN, one IPMI LAN, one VGA and two USB 3.0 ports
Front I/O: power button and UID button
Includes 27" to 32" Quick-install Rackmount Rail Kit

Please note that the 35.2" depth of this chassis (894mm) requires a rackmount cabinet with 39.3" mounting depth

(2) Intel Xeon 6126 "Skylake" 2.60 GHz 12-Core 14nm CPU - 125W TDP
with two AVX-512 units per core, 19.25MB L3 Cache, up to 768GB DDR4-2666 memory, up to three 10.4GT/s UPI links
Supports Hyper-Threading and Turbo Boost up to 3.70 GHz (clock speeds with AVX-512 instructions: 1.7-3.5 GHz)
(optional F-SKU with support for 100GBps Intel Omni-Path fabric interconnect)

(12) 64GB DDR4 ECC LRDIMM 2666MHz Memory

(768GB Total Memory @ 2666MHz)

(2) NVIDIA Tesla V100 SXM2 16GB "Volta" GPU Accelerator
SXM2 form factor with 150GB/s NVLink 2.0 interconnect (300GB/s bidirectional)
GV100 GPU chip with NVIDIA Passive Heatsink
5,120 CUDA Cores with Enhanced Unified Memory and Cooperative Groups
640 NVIDIA Tensor Cores optimized for Deep Learning training
GPU-boost capability allows for increased clock speeds
16GB High-Bandwidth HBM2 Memory (900 GB/sec peak bandwidth)
Supports INT8, INT32 integer; IEEE Half-, Single-, and Double-Precision Floating Point operations
Performance (with GPU Boost): 125 TFLOPS (half), 15.7 TFLOPS (single), 7.8 TFLOPS (double)
Power Consumption: 300W TDP
Software Development Tools:
OpenACC, OpenCL, C/C++ language compiler, debugger, profiler and memory analyzer
Standard numerical libraries: cuDNN, nvGRAPH, FFT, BLAS, SPARSE, RAND & more:
https://developer.nvidia.com/gpu-accelerated-libraries

(2) 240GB Intel DC S4500 2.5" SATA 6Gbps 3D NAND SSD
SATA 6Gb/s Interface (Supports 3Gb/s)
3D NAND TLC Internal Solid State Drive
Targeted Lifetime Endurance: ~1 Drive Write Per Day; 0.62 PBW
Full data path and Power loss protection; 256 bit AES encryption
Sustained sequential read: up to 500 MB/s
Sustained sequential write: up to 190 MB/s
Random 4KB IOPS: up to 69,000 read; up to 16,000 write
Average Latency: 36µs read, 36 µs write
2,000,000 Hours MTBF
Uncorrectable Bit Error Rate (UBER): 1 sector per 10^17 bits read

(Linux Software RAID1 Mirror Recommended; ~240GB Usable)

Qty	Description	Unit Price	Extended Price

Intel DC P4600 Series 2TB PCI-E NVMe SSD (Half-Height Half-Length)
PCI-E x4 3.1 NVMExpress Interface to System
3D NAND TLC Internal Solid State Drive
Targeted Lifetime Endurance: ~3 Drive Writes Per Day, 11.08 PBW
Built in power protection
Sequential Read: up to 3,200 MB/s
Sequential Write: up to 1,575 MB/s
Random Read IOPS: Up to 610,000
Random Write IOPS: Up to 196,650
Read Latency: 85µs
Write Latency: 15µs
2,000,000 Hours MTBF
Uncorrectable Bit Error Rate: 1 sector per 10^17 bits read

(Local Cache)

Mellanox ConnectX-5 VPI Dual-Port QSFP28 HCA, EDR InfiniBand (100Gb/s) and 100GigE
PCI-E 3.0 x16 8GT/s; Supports RDMA, SR-IOV, and multiple offload capabilities
http://www.mellanox.com/related-docs/prod_adapter_cards/PB_ConnectX-5_VPI_Card.pdf

LG Slim External USB 2.0 8X DVD/ 24X CD Writer (Black)
DVD+RW/-RW: 8X/6X
Also reads and writes DVD-RW/+RW/-RAM, CD-R/-RW
Buffer Memory: 1.5MB
Sustained Rate: DVD-ROM (max.) 11.08MB/s, CD-ROM 3.6MB/s

CentOS 7.x Linux (or your distribution of choice) & NVIDIA CUDA 9.x installed, configured and tested.

IPMI management capability is integrated on the motherboard. IPMI allows administrators to remotely monitor and control each computer individually. Capabilities include:
* System Power On, Power Off and Reset
* Monitor Fan Speeds, Component Temperatures and Voltages
* Remote Serial Console and KVM Access
* Virtual Media: local devices (CD, USB, etc) appear as directly connected to remote system

2	**NVIDIA Educational Discount**	-$1,500	-$3,000

$1,500 EDU discount per NVIDIA Tesla V100 16GB SXM2 GPU (restrictions apply)

MWYQ23427

1	**Microway Xeon + Tesla GPU 1U Compute Nodes**	$26,153	$26,153

NumberSmasher Dual Intel Xeon 1U GPU Server with
1600W 80 PLUS Platinum Redundant 1+1 Power Supplies
(power supplies require 208V power for redundancy under full load)
Up to two Intel Xeon E5-2600 Socket R3 processors
Intel C612 chipset
Intel QuickPath Interconnect (QPI) with system bus up to 9.6GT/s
Sixteen slots for up to 1024GB ECC DDR4-2400/2133 memory
Dual Integrated Intel i350 Gigabit Ethernet ports
Three PCI-E x16 3.0 slots for double-width GPUs
One PCI-E x8 3.0 Slot (Physical x16, low-profile)
Integrated AST2400 Graphics Controller
Integrated SATA3 6Gbps Controller
IPMI 2.0 w/ Virtual Media, KVM and Dedicated LAN Support
1 VGA, 2 Gigabit LAN, 1 IPMI LAN, 2 USB 3.0 ports, 1 serial header
Four hot-swap 2.5" SAS/SATA Hard Drive Bays
Also supports coprocessors
Rackmount Rail Kit Included

Qty	Description	Unit Price	Extended Price

NVIDIA Tesla P100 PCI-E 16GB "Pascal" GPU Accelerator
Full-Speed PCI-E 3.0 x16 Link
GP100 GPU chip with NVIDIA-certified Passive Heatsink
3584 CUDA Cores with Unified Memory and Page Migration Engine
GPU-boost capability allows for increased clock speeds
16GB High-Bandwidth HBM2 Memory (732 GB/sec peak bandwidth)
IEEE Half-, Single-, and Double-Precision Floating Point
Performance (with GPU Boost): 18.7 TFLOPS (half), 9.3 TFLOPS (single), 4.7 TFLOPS (double)
Double PCI slot form factor
Power Consumption: 250W TDP (One 8- and one 6-pin connector on rear)
Software Development Tools:
C/C++ language compiler, debugger, profiler and memory analyzer
Standard numerical libraries: cuDNN, nvGRAPH, FFT, BLAS, SPARSE, RAND & more:
https://developer.nvidia.com/gpu-accelerated-libraries

(2) Intel Xeon E5-2699v4 Broadwell-EP 2.20 GHz Twenty-Two Core 14nm CPU
with 55MB L3 Cache, DDR4-2400, 9.6 GT/sec QPI, 145W
Supports Hyper-Threading and Turbo Boost up to 3.6 GHz

(16) 32GB DDR4 2400 MHz ECC/Registered Memory (Dual Rank)

(512GB Total Memory @ 2133MHz)

(2) 240GB Intel DC S4500 2.5" SATA 6Gbps 3D NAND SSD
SATA 6Gb/s Interface (Supports 3Gb/s)
3D NAND TLC Internal Solid State Drive
Targeted Lifetime Endurance: ~1 Drive Write Per Day; 0.62 PBW
Full data path and Power loss protection; 256 bit AES encryption
Sustained sequential read: up to 500 MB/s
Sustained sequential write: up to 190 MB/s
Random 4KB IOPS: up to 69,000 read; up to 16,000 write
Average Latency: 36µs read, 36 µs write
2,000,000 Hours MTBF
Uncorrectable Bit Error Rate (UBER): 1 sector per 10^17 bits read

(Linux Software RAID1; Operating System)

Intel DC P4600 Series 2TB PCI-E NVMe SSD (Half-Height Half-Length)
PCI-E x4 3.1 NVMExpress Interface to System
3D NAND TLC Internal Solid State Drive
Targeted Lifetime Endurance: ~3 Drive Writes Per Day, 11.08 PBW
Built in power protection
Sequential Read: up to 3,200 MB/s
Sequential Write: up to 1,575 MB/s
Random Read IOPS: Up to 610,000
Random Write IOPS: Up to 196,650
Read Latency: 85µs
Write Latency: 15µs
2,000,000 Hours MTBF
Uncorrectable Bit Error Rate: 1 sector per 10^17 bits read

Mellanox ConnectX-5 VPI Dual-Port QSFP28 HCA, EDR InfiniBand (100Gb/s) and 100GigE
PCI-E 3.0 x16 8GT/s; Supports RDMA, SR-IOV, and multiple offload capabilities
http://www.mellanox.com/related-docs/prod_adapter_cards/PB_ConnectX-5_VPI_Card.pdf

CentOS 7.x Linux (or your distribution of choice) & NVIDIA CUDA 9.x installed, configured and tested.

IPMI management capability is integrated on the motherboard. IPMI allows administrators to remotely monitor and control each computer individually. Capabilities include:
* System Power On, Power Off and Reset
* Monitor Fan Speeds, Component Temperatures and Voltages
* Remote Serial Console and KVM Access
* Virtual Media: local devices (CD, USB, etc) appear as directly connected to remote system

Qty	Description	Unit Price	Extended Price
1	**NVIDIA Academic Instant Rebate** $1,500 EDU discount per NVIDIA Tesla P100 16GB GPU (restrictions apply)	-$1,500	-$1,500

MWYQ23428

Qty	Description	Unit Price	Extended Price
1	**Microway Xeon + Tesla GPU 1U Compute Nodes**	$17,301	$17,301

NumberSmasher Dual Intel Xeon 1U GPU Server with
1600W 80 PLUS Platinum Redundant 1+1 Power Supplies
(power supplies require 208V power for redundancy under full load)
Up to two Intel Xeon E5-2600 Socket R3 processors
Intel C612 chipset
Intel QuickPath Interconnect (QPI) with system bus up to 9.6GT/s
Sixteen slots for up to 1024GB ECC DDR4-2400/2133 memory
Dual Integrated Intel i350 Gigabit Ethernet ports
Three PCI-E x16 3.0 slots for double-width GPUs
One PCI-E x8 3.0 Slot (Physical x16, low-profile)
Integrated AST2400 Graphics Controller
Integrated SATA3 6Gbps Controller
IPMI 2.0 w/ Virtual Media, KVM and Dedicated LAN Support
1 VGA, 2 Gigabit LAN, 1 IPMI LAN, 2 USB 3.0 ports, 1 serial header
Four hot-swap 2.5" SAS/SATA Hard Drive Bays
Also supports coprocessors
Rackmount Rail Kit Included

NVIDIA Tesla P100 PCI-E 16GB "Pascal" GPU Accelerator
Full-Speed PCI-E 3.0 x16 Link
GP100 GPU chip with NVIDIA-certified Passive Heatsink
3584 CUDA Cores with Unified Memory and Page Migration Engine
GPU-boost capability allows for increased clock speeds
16GB High-Bandwidth HBM2 Memory (732 GB/sec peak bandwidth)
IEEE Half-, Single-, and Double-Precision Floating Point
Performance (with GPU Boost): 18.7 TFLOPS (half), 9.3 TFLOPS (single), 4.7 TFLOPS (double)
Double PCI slot form factor
Power Consumption: 250W TDP (One 8- and one 6-pin connector on rear)
Software Development Tools:
C/C++ language compiler, debugger, profiler and memory analyzer
Standard numerical libraries: cuDNN, nvGRAPH, FFT, BLAS, SPARSE, RAND & more:
https://developer.nvidia.com/gpu-accelerated-libraries

(2) Intel Xeon E5-2637v4 Broadwell-EP 3.50 GHz Four Core 14nm CPU
with 15MB L3 Cache, DDR4-2400, 9.6 GT/sec QPI, 135W
Supports Hyper-Threading and Turbo Boost up to 3.7 GHz

(16) 32GB DDR4 2400 MHz ECC/Registered Memory (Dual Rank)

(256GB Total Memory @ 2400MHz)

240GB Intel DC S4500 2.5" SATA 6Gbps 3D NAND SSD
SATA 6Gb/s Interface (Supports 3Gb/s)
3D NAND TLC Internal Solid State Drive
Targeted Lifetime Endurance: ~1 Drive Write Per Day; 0.62 PBW
Full data path and Power loss protection; 256 bit AES encryption
Sustained sequential read: up to 500 MB/s
Sustained sequential write: up to 190 MB/s
Random 4KB IOPS: up to 69,000 read; up to 16,000 write
Average Latency: 36µs read, 36 µs write
2,000,000 Hours MTBF
Uncorrectable Bit Error Rate (UBER): 1 sector per 10^17 bits read

Mellanox ConnectX-5 VPI Dual-Port QSFP28 HCA, EDR InfiniBand (100Gb/s) and 100GigE
PCI-E 3.0 x16 8GT/s; Supports RDMA, SR-IOV, and multiple offload capabilities
http://www.mellanox.com/related-docs/prod_adapter_cards/PB_ConnectX-5_VPI_Card.pdf

CentOS 7.x Linux (or your distribution of choice) & NVIDIA CUDA 9.x installed, configured and tested.

Qty	Description	Unit Price	Extended Price

IPMI management capability is integrated on the motherboard. IPMI allows administrators to remotely monitor and control each computer individually. Capabilities include:
* System Power On, Power Off and Reset
* Monitor Fan Speeds, Component Temperatures and Voltages
* Remote Serial Console and KVM Access
* Virtual Media: local devices (CD, USB, etc) appear as directly connected to remote system

1	**NVIDIA Academic Instant Rebate**	-$1,500	-$1,500

$1,500 EDU discount per NVIDIA Tesla P100 16GB GPU (restrictions apply)

MWYQ23434

1	**Microway 2U EPYC Compute Server**	$29,024	$29,024

Navion Dual AMD EPYC 2U Server
1600W High-Efficiency "Titanium" Redundant 1+1 Power Supply
Up to two AMD EPYC 7000-Series processors (Socket SP3 LGA)
AMD Infinity Fabric link between CPU sockets
Thirty-two slots for up to 4TB ECC DDR4-2666 memory
Quad Integrated Intel i350AM4 Gigabit Ethernet Ports (RJ45)
Integrated ASPEED AST2500 BMC with Graphics Controller
IPMI 2.0 w/ Virtual Media, KVM and Dedicated LAN Support
PCI-Express expansion slots:
One full-height PCI-Express 3.0 x16 slot (10.5" length)
Five full-height PCI-Express 3.0 x8 slots (9.5" length)
One low-profile PCI-Express 3.0 x8 slot
One low-profile PCI-Express 3.0 x8 (internal)
Twelve front hot-swap 3.5" drive bays supporting either:
Default: (12) SATA 6Gbps drives (via onboard controller)
(8) SATA 6Gbps drives and (4) NVMe drives (with optional add-in controller)
(SATA 6Gbps bays can be upgraded to SATA/SAS 12Gbps bays via optional add-in controllers)
Two rear hot-swap 2.5" drive bays (optional)
Internal ports: two SATA SuperDOM, one TPM header, and one USB 3.0 port (Type A)
Rear ports: four 1G LAN, one IPMI LAN, one VGA, two USB 3.0, and one serial COM port
Front I/O: power button and reset button
Rackmount Rail Kit included (29" chassis depth)

(2) AMD "EPYC" 7551 Thirty Two Core CPU 2.0 GHz Base Clock Frequency
2.55 GHz All Core Boost Speed, 3.0 GHz Max Core Boost Speed
64MB Shared L3 Cache, 180W TDP

(16) 32GB DDR4 2666 MHz ECC/Registered Memory (Dual Rank)

(512GB Total Memory @ 2666MHz)

240GB Intel DC S4500 2.5" SATA 6Gbps 3D NAND SSD
SATA 6Gb/s Interface (Supports 3Gb/s)
3D NAND TLC Internal Solid State Drive
Targeted Lifetime Endurance: ~1 Drive Write Per Day; 0.62 PBW
Full data path and Power loss protection; 256 bit AES encryption
Sustained sequential read: up to 500 MB/s
Sustained sequential write: up to 190 MB/s
Random 4KB IOPS: up to 69,000 read; up to 16,000 write
Average Latency: 36µs read, 36 µs write
2,000,000 Hours MTBF
Uncorrectable Bit Error Rate (UBER): 1 sector per 10^17 bits read

(Operating System)

2.5" HDD Adapter for 3.5" HDD Carrier Tray (Tool-less)

Qty	Description	Unit Price	Extended Price

Intel DC P4600 Series 2TB PCI-E NVMe SSD (Half-Height Half-Length)
PCI-E x4 3.1 NVMExpress Interface to System
3D NAND TLC Internal Solid State Drive
Targeted Lifetime Endurance: ~3 Drive Writes Per Day, 11.08 PBW
Built in power protection
Sequential Read: up to 3,200 MB/s
Sequential Write: up to 1,575 MB/s
Random Read IOPS: Up to 610,000
Random Write IOPS: Up to 196,650
Read Latency: 85µs
Write Latency: 15µs
2,000,000 Hours MTBF
Uncorrectable Bit Error Rate: 1 sector per 10^17 bits read

(Local Scratch)

Mellanox ConnectX-5 VPI Dual-Port QSFP28 HCA, EDR InfiniBand (100Gb/s) and 100GigE
PCI-E 3.0 x16 8GT/s; Supports RDMA, SR-IOV, and multiple offload capabilities
http://www.mellanox.com/related-docs/prod_adapter_cards/PB_ConnectX-5_VPI_Card.pdf

NVIDIA Tesla V100 PCI-E 16GB "Volta" GPU Accelerator
Full-Speed PCI-E 3.0 x16 Link
GV100 GPU chip with NVIDIA Passive Heatsink
5,120 CUDA Cores with Enhanced Unified Memory and Cooperative Groups
640 NVIDIA Tensor Cores optimized for Deep Learning training
GPU-boost capability allows for increased clock speeds
16GB High-Bandwidth HBM2 Memory (900 GB/sec peak bandwidth)
Supports INT8, INT32 integer; IEEE Half-, Single-, and Double-Precision Floating Point operations
Performance (with GPU Boost): 112 TFLOPS (half), 14 TFLOPS (single), 7 TFLOPS (double)
Double PCI slot form factor
Power Consumption: 250W TDP (One 8- and one 6-pin connector on rear)
Software Development Tools:
OpenACC, OpenCL, C/C++ language compiler, debugger, profiler and memory analyzer
Standard numerical libraries: cuDNN, nvGRAPH, FFT, BLAS, SPARSE, RAND & more:
https://developer.nvidia.com/gpu-accelerated-libraries

LG Slim External USB 2.0 8X DVD/ 24X CD Writer (Black)
DVD+RW/-RW: 8X/6X
Also reads and writes DVD-RW/+RW/-RAM, CD-R/-RW
Buffer Memory: 1.5MB
Sustained Rate: DVD-ROM (max.) 11.08MB/s, CD-ROM 3.6MB/s

CentOS 7.x Linux (or your distribution of choice) & NVIDIA CUDA 9.x installed, configured and tested.

IPMI management capability is integrated on the motherboard. IPMI allows administrators to remotely monitor and control each computer individually. Capabilities include:
* System Power On, Power Off and Reset
* Monitor Fan Speeds, Component Temperatures and Voltages
* Remote Serial Console and KVM Access
* Virtual Media: local devices (CD, USB, etc) appear as directly connected to remote system

1	**NVIDIA Academic Instant Rebate**	-$1,500	-$1,500

$1,500 EDU discount per NVIDIA Tesla V100 16GB GPU (restrictions apply)

MWYQ23484

Qty	Description	Unit Price	Extended Price
1	**Mellanox 40/10Gb Ethenet Network** (2) Mellanox SN2700B 32-Port 40GbE Ethernet Managed Switch (1U) 6.4Tb/s linerate performance (up to 9.52 billion packets-per-second) 32 front-mounted QSFP28 ports (40/10GbE) Supports 40GbE to 4 x 10GbE via breakout cables 1 management port, 1 serial port Dual, Redundant Hot-Swap Power Supplies Standard depth, Power supply side to connector side airflow Rackmount Rail Kit Included Model MSN2700-BS2F (includes 2 year Bronze warranty) (32) 40GigE Mellanox Ethernet Cable, QSFP+ Passive Copper, 2 meters, 30 AWG	$29,102	$29,102
	MWYQ23493-01		
1	**Cluster Management Licenses** (20) Bright Cluster Manager (Advanced Edition - Academic) Advanced, Easy-to-Use Cluster Management Software Subscription License (priced per year) 1-year standard support and software maintenance subscription, including access to the Bright Computing RPM repository. http://www.brightcomputing.com/product-offerings/bright-cluster-manager-for-hpc/	$2,880	$2,880
	Shipping and Insurance: prepay and add		
	Total (Academic Pricing)		**$183,487**

FOB: Destination

Ship Date: 4 - 5 Weeks after receipt of order and documentation

Warranty: DGX Provided through NVIDIA; All other items are two years offsite with replacement components typically cross-shipped within 24 hours of problem determination by Microway Tech Support.

Technical Support: Lifetime technical support via telephone, fax, or email.

Terms: NET 30

Shipping Method: FedEx Freight

The Buyer is responsible for any sales taxes or duties related to the purchase.

NVIDIA requires the purchase of a 1-Year or 3-Year Support Contract

All DGX orders are non-cancellable and non-refundable.

Since 1982 Microway has been a leader in providing high performance computing solutions. Microway specializes in building complex clusters, servers, and workstations. We are unique in having Linux expertise throughout our organization to provide testing of all systems at our assembly and integration center at our headquarters in Plymouth, Massachusetts. Our validation suite includes a number of MPI applications and Microway proprietary software, including MPI Link-Checker and InfiniScope.

Microway is classified as a small business - woman owned and operated.

Microway welcomes our customers (and potential customers) to personally visit our manufacturing facility. We value the opportunity to share our understanding of the systems we build, and to demonstrate our dedication to quality in our design, fabrication, final testing and technical support. Please contact me if you plan to be in or near Massachusetts and would like to make an appointment.

Prices subject to review at time of order due to world-wide component price volatility.

This quote is valid for 90 days.